

COP4610: Operating Systems

Project 1

Zhi Wang

Florida State University

Spring 2015

Run xv6

- `ssh linprog.cs.fsu.edu`
- `wget http://www.cs.fsu.edu/~zwang/files/cop4610/Spring2015/xv6_v8.tar.gz`
- `tar -xf xv6_v8.tar.gz`
- `cd xv6`
- to compile only: `make`
 - ▮ to compile and run: `make qemu-nox`
- to quit qemu: `ctrl-a x`

User Space

- We use `close` as the example, `read` is similar.
 - ➡ assuming a program calls `close (5)`, `5` is an invalid file descriptor

User Space

- We use `close` as the example, `read` is similar.
 - ▮ assuming a program calls `close (5)`, `5` is an invalid file descriptor
- `close` is declared in `user.h`: `int close(int);`
- `close` is defined in `usys.S`: `SYSCALL(close)`, which expands to
 - `.globl close;` → declare `close` as a global symbol
 - `close:` → definition of `close`
 - `movl $SYS_close, %eax;` → put system call number in register `eax`
 - `int $T_SYSCALL;` → trigger a software interrupt, enter the kernel
 - `ret` → return to the caller of `close`

User Space

- We use `close` as the example, `read` is similar.
 - ▮ assuming a program calls `close (5)`, `5` is an invalid file descriptor
- `close` is declared in `user.h`: `int close(int);`
- `close` is defined in `usys.S`: `SYSCALL(close)`, which expands to
 - `.globl close; →declare close as a global symbol`
 - `close: →definition of close`
 - `movl $SYS_close, %eax; →put system call number in register eax`
 - `int $T_SYSCALL; →trigger a software interrupt, enter the kernel`
 - `ret →return to the caller of close`
- `T_SYSCALL` defined in `traps.h`, `SYS_close` defined in `syscall.h`

Entering the Kernel

- `int $T_SYSCALL` triggers a software interrupt (`T_SYSCALL=64`)
 - ➡ CPU saves the current state, and calls the interrupt handler
 - ➡ Interrupt handler for `T_SYSCALL` is `vector64` (`vectors.S`)
 - ➡ `vector64` jumps to `alltraps` function, which creates the `trapframe`, and calls `trap` (`struct trapframe *tf`)

Entering the Kernel

- `int $T_SYSCALL` triggers a software interrupt (`T_SYSCALL=64`)
 - ⇒ CPU saves the current state, and calls the interrupt handler
 - ⇒ Interrupt handler for `T_SYSCALL` is `vector64` (`vectors.S`)
 - ⇒ `vector64` jumps to `alltraps` function, which creates the `trapframe`, and calls `trap (struct trapframe *tf)`
- `struct trapframe` saves the user-space registers. `tf→eax` contains the system call number (`SYS_close`)

```
struct trapframe {
    uint edi;
    uint esi;
    uint ebp;
    uint oesp;
    uint ebx;
    uint edx;
    ...
}
```

Syscall Dispatch

- `trap (tf)` calls `syscall (void)` because `(tf→trapno == T_SYSCALL)`
 - ▮▮▮ `trapframe tf` saved to the current process control block
- `syscall` reads the syscall number in `eax`, and calls `sys_close`
 - ▮▮▮ `syscalls[SYS_close]`
 - ▮▮▮ return value is saved in `tf→eax`, the kernel restores `tf` before returning to user space

sys_close and Return

- `sys_close` reads the parameter from user stack with `argfd`
 - ▮ `fd = 5`, an invalid file descriptor
 - ▮ `sys_close` returns `-1`
- it returns to `syscall(void)`, which saves return value to `eax` and returns to `trap`
- `trap` returns to `alltraps`, which restores user registers and returns to user space with `iret`

Define a New Syscall

- User space:
 - ➡ declare `getprocs` in `user.h`, define `getprocs` in `usys.S`
 - ➡ create a program called `ps` and add it to the makefile (refer `cat.c`)
- Kernel space:
 - ➡ add a new system call number: `SYS_ps` in `syscall.h`
 - ➡ create the syscall handler (`sys_ps`) in `sysproc.c` for `SYS_ps`
 - ➡ add `sys_ps` to the syscall dispatcher `syscalls`

Hint:

follow `close/sys_close` as an example to complete this part.

sys_ps

- `proc.c` has a global variable `ptable` with all the process control blocks
- `sys_ps` copies each PCB to user-provided memory
 - ➡ note: not all PCB states are copied
 - ➡ check the size of the user-provided memory